

Name: Nour Mohamed Mahmoud

First we make importing required libraries, and make a classification of USA Cars dataset.

```
jupyter jupyter Last Checkpoint: 48 minutes ago
File Edit View Run Kernel Settings Help Trusted
JupyterLab Python 3 (ipykernel)
```

```
[1]: import pandas
dataframe = pandas.read_csv('mushrooms.csv')
dataframe

[1]:
```

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	stalk-surface-below-ring	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	ring-number	ring-type	spore-print-color	population	hazards
0	p	x	s	n	t	p	f	c	n	k	...	s	w	w	p	w	o	p	k	s
1	e	x	s	y	t	a	f	c	b	k	...	s	w	w	p	w	o	p	n	n
2	e	b	s	w	t	l	f	c	b	n	...	s	w	w	p	w	o	p	n	n
3	p	x	y	w	t	p	f	c	n	n	...	s	w	w	p	w	o	p	k	s
4	e	x	s	g	f	n	f	w	b	k	...	s	w	w	p	w	o	e	n	a
...
8119	e	k	s	n	f	n	a	c	b	y	...	s	o	o	p	o	o	p	b	c
8120	e	x	s	n	f	n	a	c	b	y	...	s	o	o	p	n	o	p	b	v
8121	e	f	s	n	f	n	a	c	b	n	...	s	o	o	p	o	o	p	b	c
8122	p	k	y	n	f	y	f	c	n	b	...	k	w	w	p	w	o	e	w	v
8123	e	x	s	n	f	n	a	c	b	y	...	s	o	o	p	o	o	p	o	c

8124 rows x 23 columns

Importing the required libraries to show the dataset.

```
] : import numpy as np
import pandas as pd
import os
for dirname, _, filenames in os.walk('/jupyter/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
: dataset=pd.read_csv("USA_Cars.csv")
```

```
dataset.head()
```

```
: Unnamed: 0 price brand model year title_status mileage color vin lot state country condition
0 0 6300 toyota cruiser 2008 clean vehicle 274117.0 black jtezu11f88k007763 159348797 new jersey usa 10 days left
1 1 2899 ford se 2011 clean vehicle 190552.0 silver 2fmdk3gc4bbb02217 166951262 tennessee usa 6 days left
2 2 5350 dodge mpv 2018 clean vehicle 39590.0 silver 3c4pdcgg5jt346413 167655728 georgia usa 2 days left
3 3 25000 ford door 2014 clean vehicle 64146.0 blue 1ftfw1et4efc23745 167753855 virginia usa 22 hours left
4 4 27700 chevrolet 1500 2018 clean vehicle 6654.0 red 3gcpcrec2jg473991 167763266 florida usa 22 hours left
```

```
[6]: dataset.shape
```

```
[6]: (2499, 13)
```

```
3]: dataset.duplicated().sum()
```

```
-  
[8]: 0
```

```
[11]: dataset=dataset.drop(["Unnamed: 0", "vin", "lot", "condition"],axis=1)
```

```
[12]: dataset.describe()
```

```
[12]:
```

	price	year	mileage
count	2499.000000	2499.000000	2.499000e+03
mean	18767.671469	2016.714286	5.229869e+04
std	12116.094936	3.442656	5.970552e+04
min	0.000000	1973.000000	0.000000e+00
25%	10200.000000	2016.000000	2.146650e+04
50%	16900.000000	2018.000000	3.536500e+04
75%	25555.500000	2019.000000	6.347250e+04
max	84900.000000	2020.000000	1.017936e+06

```
[13]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 2499 entries, 0 to 2498  
Data columns (total 9 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   price           2499 non-null   int64  
1   brand           2499 non-null   object  
2   model           2499 non-null   object  
3   year            2499 non-null   int64  
4   title_status    2499 non-null   object  
5   mileage         2499 non-null   float64  
6   color           2499 non-null   object  
7   state           2499 non-null   object  
8   country         2499 non-null   object  
dtypes: float64(1), int64(2), object(6)  
memory usage: 175.8+ KB
```

```
[14]: col1=["brand","model","title_status","color","state","country"]
      for col in col1:
          dataset[col]=dataset[col].str.strip()
```

```
[15]: count=dataset["brand"].value_counts()
      count
```

```
[15]: brand
      ford          1235
      dodge          432
      nissan          312
      chevrolet      297
      gmc             42
      jeep           30
      chrysler       18
      bmw            17
      hyundai        15
      buick          13
      kia            13
      infiniti       12
      honda          12
      mercedes-benz  10
      cadillac       10
      heartland      5
      audi           4
      land           4
      peterbilt      4
      acura          3
      mazda          2
      lexus          2
      lincoln        2
      toyota         1
      harley-davidson 1
      jaguar         1
      maserati       1
      ram            1
      Name: count, dtype: int64
```

```
[16]: threshold=9
      rep=count[count<=threshold].index
      dataset["brand"]=dataset["brand"].replace(rep,"uncommon")
```

```
[18]: dataset["brand"].value_counts()
```

```
[18]: brand
      ford          1235
      dodge         432
      nissan         312
      chevrolet     297
      gmc           42
      uncommon      31
      jeep          30
      chrysler      18
      bmw           17
      hyundai       15
      buick         13
      kia           13
      honda         12
      infiniti      12
      mercedes-benz 10
      cadillac      10
      Name: count, dtype: int64
```

```
[19]: count1=dataset["model"].value_counts()
```

```
[20]: threshold=10
      rep=count1[count1<=threshold].index
```

```
[20]: threshold=10
      rep=count1[count1<=threshold].index
      dataset["model"]=dataset["model"].replace(rep,"other_models")
```

```
[21]: dataset["model"].value_counts()
```

```

[21]: model
      door          651
      other_models  255
      f-150         219
      doors         148
      caravan       102
      mpv           87
      fusion        65
      durango       64
      journey       61
      rogue         54
      van           46
      challenger    44
      charger       42
      transit       41
      max           41
      sport         40
      1500          39
      escape        39
      explorer      39
      srw           38
      versa         34
      edge          34
      flex          33
      wagon         30
      mustang       29
      sentra        28
      expedition    28
      pathfinder    22
      altima        21
      equinox       18
      suburban      16
      pickup        15
      frontier      14

fiesta          14
malibu          12
colorado        12
cutaway         12
impala          12
Name: count, dtype: int64

```

```
[22]: dataset["model"].nunique()
```

```
[22]: 38
```

```
[23]: dataset["title_status"].value_counts()
```

```
[23]: title_status  
clean vehicle      2336  
salvage insurance   163  
Name: count, dtype: int64
```

```
[25]: count2=dataset["color"].value_counts()  
count2
```

```
[25]: color  
white      707  
black      516  
gray       395  
silver     300  
red        192  
blue       151  
no_color   61  
green      24  
orange     20  
gold       19  
charcoal   18  
brown      15  
yellow     9  
magnetic metallic 6  
shadow black 5  
color:      5  
beige       5  
ingot silver metallic 4  
oxford white 4  
billet silver metallic clearcoat 3  
triple yellow tri-coat 3  
super black 3  
off-white 2  
ruby red metallic tinted clearcoat 2  
cayenne red 2  
white platinum tri-coat metallic 2  
tuxedo black metallic 2  
black clearcoat 2  
bright white clearcoat 2  
phantom black 1  
maroon      1  
dark blue   1  
turquoise   1  
purple      1
```

competition orange	1
toreador red	1
jazz blue pearlcoat	1
light blue	1
kona blue metallic	1
royal crimson metallic tinted clearcoat	1
ruby red	1
guard	1
ingot silver	1
lightning blue	1
tan	1
burgundy	1
morningsky blue	1
pearl white	1
glacier white	1

Name: count, dtype: int64

```
8]: threshold=14
    rep2=count2[count2<=threshold].index
    dataset["color"]=dataset["color"].replace(rep2,"other_colors")
```

```
[29]: dataset["color"].value_counts()
```

```
[29]: color
      white      707
      black      516
      gray       395
      silver     300
      red        192
      blue       151
      other_colors  81
      no_color    61
      green       24
      orange      20
      gold        19
      charcoal    18
      brown       15
      Name: count, dtype: int64
```

```
[31]: count3=dataset["state"].value_counts()
```

```
[32]: dataset["state"].nunique()
```

```
[32]: 44
```

```
[33]: threshold=50  
rep3=count3[count3<=threshold].index  
dataset["state"]=dataset["state"].replace(rep3,"other states")
```

```
[33]: threshold=50  
rep3=count3[count3<=threshold].index  
dataset["state"]=dataset["state"].replace(rep3,"other_states")
```

```
[34]: dataset["state"].value_counts()
```

```
[34]: state  
other_states      403  
pennsylvania      299  
florida           246  
texas             214  
california        190  
michigan          169  
north carolina    146  
minnesota         119  
illinois          113  
wisconsin         94  
virginia          90  
new jersey        87  
nevada            85  
oklahoma          71  
south carolina    64  
new york          58  
georgia           51  
Name: count, dtype: int64
```

```
[35]: dataset["country"].value_counts()
```



```
[35]: country
      usa      2492
      canada    7
      Name: count, dtype: int64
```

```
[36]: col=["brand","model","year","title_status","mileage","color","state","country"]
      X_train=dataset[col]
      Y_train=dataset["price"]
```

```
[37]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(X_train,Y_train,test_size=0.3)
```

```
[38]: from sklearn.preprocessing import OneHotEncoder
      from sklearn.preprocessing import OrdinalEncoder
      from sklearn.compose import ColumnTransformer

      ct=ColumnTransformer([("tnf1",OrdinalEncoder(categories=[["clean vehicle","salvage insurance"]]),["title_status"]),
                             ("tnf2",OneHotEncoder(sparse_output=False,drop="first"),["brand","model","color","state","country",])],remainder="passthrough")
```

```
•[39]: x_train_column_transform=ct.fit_transform(x_train)
      x_test_column_transform=ct.transform(x_test)
```

```
[40]: x_train_column_transform.shape
```

```
[40]: (1749, 84)
```

```
[41]: from sklearn.tree import DecisionTreeRegressor
      from sklearn.svm import SVR
      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_absolute_percentage_error, mean_absolute_error
      from sklearn.model_selection import train_test_split
```

```
[42]: lr=LinearRegression()
      sv=SVR()
      dt=DecisionTreeRegressor()
```

```
[43]: lr.fit(x_train_column_transform,y_train)
      sv.fit(x_train_column_transform,y_train)
      dt.fit(x_train_column_transform,y_train)
```

```
[43]: ▾ DecisionTreeRegressor  
DecisionTreeRegressor()
```

```
[44]: lr_pred=lr.predict(x_test_column_transform)  
sv_pred=sv.predict(x_test_column_transform)  
dt_pred=dt.predict(x_test_column_transform)
```

```
[45]: from sklearn.metrics import mean_squared_error  
print("mean_squared_error of LR:", mean_squared_error(y_test,lr_pred))  
print("mean_squared_error of SVM:", mean_squared_error(y_test,sv_pred))  
print("mean_squared_error of Decission Tree:", mean_squared_error(y_test,dt_pred))
```

mean_squared_error of LR: 54258268.27096521

mean_squared_error of SVM: 137661516.1647802

mean_squared_error of Decission Tree: 90425432.89733334

```
[46]: from sklearn.preprocessing import LabelEncoder  
label_encoders = {}  
for col in dataset.select_dtypes(include='object').columns:  
    if col != 'title_status':  
        le = LabelEncoder()  
        dataset[col] = le.fit_transform(dataset[col])  
        label_encoders[col] = le
```

```
[55]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

dataset = pd.read_csv('USA_Cars.csv')

label_encoders = {}
for col in dataset.select_dtypes(include='object').columns:
    if col != 'title_status':
        le = LabelEncoder()
        dataset[col] = le.fit_transform(dataset[col])
        label_encoders[col] = le

target_encoder = LabelEncoder()
dataset['title_status'] = target_encoder.fit_transform(dataset['title_status'])

X = dataset.drop('title_status', axis=1)
y = dataset['title_status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\n Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\n Classification Report:\n", classification_report(y_test, y_pred, target_names=target_encoder.classes_))
```

Accuracy: 0.986

Confusion Matrix:


```
[[464  1]
 [ 6 29]]
```

Classification Report:

	precision	recall	f1-score	support
clean vehicle	0.99	1.00	0.99	465
salvage insurance	0.97	0.83	0.89	35
accuracy			0.99	500
macro avg	0.98	0.91	0.94	500
weighted avg	0.99	0.99	0.99	500

```
[56]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)
labels = target_encoder.classes_
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.tight_layout()
plt.show()
```

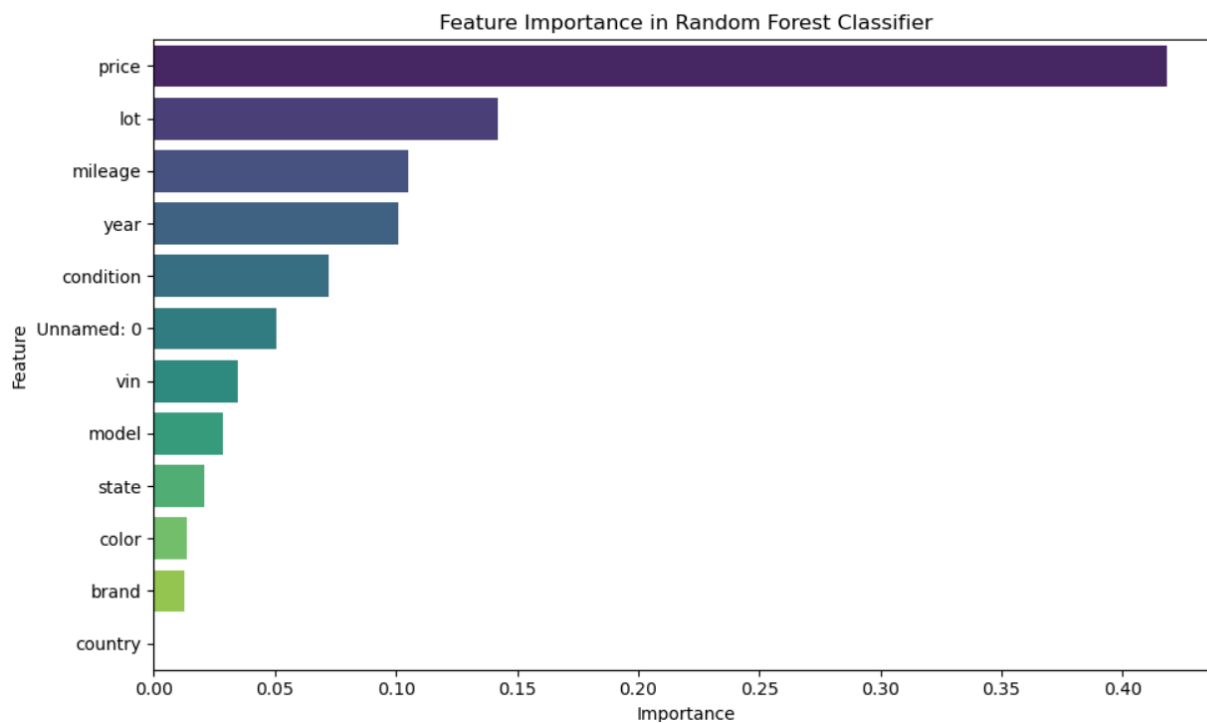


```
[58]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

importances = model.feature_importances_
feature_names = X.columns

feat_dataset = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

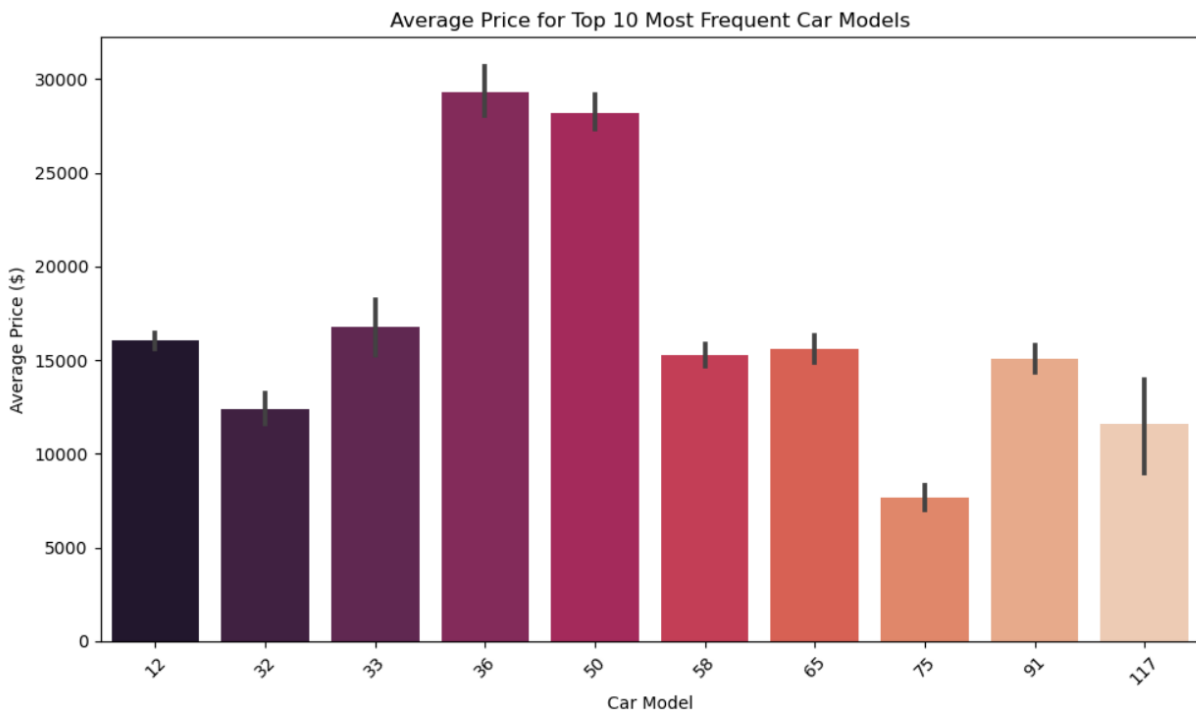
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feat_dataset, palette='viridis')
plt.title('Feature Importance in Random Forest Classifier')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()
```



```
[61]: top_models = dataset['model'].value_counts().head(10).index
      filtered_dataset = dataset[dataset['model'].isin(top_models)]

      plt.figure(figsize=(10, 6))
      sns.barplot(data=filtered_dataset, x='model', y='price', estimator='mean', palette='rocket')

      plt.title('Average Price for Top 10 Most Frequent Car Models')
      plt.xlabel('Car Model')
      plt.ylabel('Average Price ($)')
      plt.xticks(rotation=45)
      plt.tight_layout()
      plt.show()
```




```
[66]: import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns

      dataset['mileage_bin'] = pd.cut(dataset['mileage'], bins=10)

      grouped = dataset.groupby('mileage_bin')['price'].mean().reset_index()

      plt.figure(figsize=(12, 6))
      sns.barplot(data=grouped, x='mileage_bin', y='price', palette='crest')

      plt.title('Average Price per Mileage Range')
      plt.xlabel('Mileage Range')
      plt.ylabel('Average Price ($)')
      plt.xticks(rotation=45)
      plt.tight_layout()
      plt.show()
```



```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix, roc_curve, auc
import matplotlib.pyplot as plt
import seaborn as sns

# Models
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder

dataset = pd.read_csv("USA_Cars.csv")

dataset = dataset.drop(columns=["Unnamed: 0", "vin", "lot", "country"])

label_encoders = {}
for col in dataset.select_dtypes(include="object").columns:
    le = LabelEncoder()
    dataset[col] = le.fit_transform(dataset[col])
    label_encoders[col] = le

x = dataset.drop(columns=["price"])
y = dataset["price"]

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
x_train_scaled = pd.DataFrame(scaler.fit_transform(x_train), columns=x.columns)
x_test_scaled = pd.DataFrame(scaler.transform(x_test), columns=x.columns)
```

```

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, r2_score

lr = LinearRegression()
lr.fit(x_train_scaled, y_train)
y_pred = lr.predict(x_test_scaled)

print("MAE:", mean_absolute_error(y_test, y_pred))
print("R²:", r2_score(y_test, y_pred))

```

MAE: 7780.950477956126

R²: 0.2969808624192337

```

x_train_scaled = np.nan_to_num(x_train_scaled)
x_test_scaled = np.nan_to_num(x_test_scaled)

```

```

print("size of X_train_scaled:", x_train_scaled.shape)
print("size of X_test_scaled:", x_test_scaled.shape)

```

size of X_train_scaled: (1999, 8)

size of X_test_scaled: (500, 8)

```
print(y_train.value_counts())
```

```

price
0      34
16500   21
15000   18
25      18
25000   16
..
18497    1
9210     1
4450     1
1400     1
4650     1
Name: count, Length: 693, dtype: int64

```



```
models = {  
    "Logistic Regression": LogisticRegression(),  
    "Decision Tree": DecisionTreeClassifier(),  
    "SVM": SVC(probability=True),  
    "Random Forest": RandomForestClassifier(),  
    "KNN": KNeighborsClassifier()  
}  
  
results = {}  
for name, model in models.items():  
    model.fit(x_train_scaled, y_train)  
    y_pred = model.predict(x_test_scaled)  
    acc = accuracy_score(y_test, y_pred)  
    results[name] = acc  
    print(f"{name} Accuracy: {acc:.4f}")
```

Logistic Regression Accuracy: 0.0380

Decision Tree Accuracy: 0.0640

SVM Accuracy: 0.0440

Random Forest Accuracy: 0.0800

KNN Accuracy: 0.0540

```

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "KNN": KNeighborsClassifier(),
    "SVM": SVC(probability=True),
    "Random Forest": RandomForestClassifier(random_state=42)
}

results = {}
for name, model in models.items():
    model.fit(x_train_scaled, y_train)
    y_pred = model.predict(x_test_scaled)
    results[name] = model.score(x_test_scaled, y_test)

best_model_name = max(results, key=results.get)
best_model = models[best_model_name]
y_pred = best_model.predict(x_test_scaled)
y_proba = best_model.predict_proba(x_test_scaled)

print("\nBest Model:", best_model_name)
print(classification_report(y_test, y_pred))

# Confusion matrix

# طباعة مقاييس الأداء
print(f"\n✅ Best Model: {best_model_name}")
print("MAE:", mean_absolute_error(y_test, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R² Score:", r2_score(y_test, y_pred))

# ✅ Scatter Plot: Actual vs. Predicted Prices
plt.figure(figsize=(7,6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.6, color='royalblue', edgecolor='black')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title(f"Actual vs Predicted Prices ({best_model_name})", fontsize=14)
plt.xlabel("Actual Price", fontsize=12)
plt.ylabel("Predicted Price", fontsize=12)
plt.grid(True, linestyle='--', alpha=0.3)
plt.tight_layout()
plt.show()

```

```

from sklearn.preprocessing import label_binarize
from sklearn.metrics import roc_auc_score

y_test_bin = label_binarize(y_test, classes=[0, 1, 2])
fpr, tpr, roc_auc = {}, {}, {}

for i in range(3):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

plt.figure(figsize=(7,5))
for i in range(3):
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title(f"Multiclass ROC Curve ({best_model_name})")
plt.legend()
plt.tight_layout()
plt.show()

```

Best Model: Random Forest

	precision	recall	f1-score	support
0	0.47	0.78	0.58	9
25	0.00	0.00	0.00	0
50	1.00	1.00	1.00	1
75	0.00	0.00	0.00	1
accuracy			0.08	500
macro avg	0.04	0.04	0.04	500
weighted avg	0.10	0.08	0.08	500


✅ Best Model: Random Forest

MAE: 5561.428


RMSE: 8981.757954431861

R² Score: 0.4910061526938454

Actual vs Predicted Prices (Random Forest)



Multiclass ROC Curve (Random Forest)



```

import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# 1. Load and preprocess data
dataset = pd.read_csv("USA_Cars.csv")
dataset = dataset.drop(columns=["Unnamed: 0", "vin", "lot", "country"])

# Encode categorical features
for col in dataset.select_dtypes(include='object').columns:
    dataset[col] = LabelEncoder().fit_transform(dataset[col])

# Features and target
x = dataset.drop(columns=["price"])
y = dataset["price"]

# Train-test split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 2. Define parameter grid
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

# 3. Setup GridSearchCV
rf = RandomForestRegressor(random_state=42)
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid,
                           cv=3, scoring='r2', n_jobs=-1, verbose=1)

# 4. Fit the model
grid_search.fit(X_train_scaled, y_train)

# 5. Results
best_rf = grid_search.best_estimator_
print("✅ Best Parameters:", grid_search.best_params_)
print("🏆 Best Cross-Validated R² Score:", grid_search.best_score_)

# 6. Evaluation on test set
y_pred = best_rf.predict(X_test_scaled)
print("\n📊 Performance of Tuned Random Forest:")
print("MAE:", mean_absolute_error(y_test, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R² Score:", r2_score(y_test, y_pred))

```

Fitting 3 folds for each of 24 candidates, totalling 72 fits

✅ Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 200}

🏆 Best Cross-Validated R² Score: 0.6194780492869159

📊 Performance of Tuned Random Forest:

MAE: 4556.89337

RMSE: 6954.512725911367

R² Score: 0.6948430606086335

```
import seaborn as sns
import matplotlib.pyplot as plt
importances = best_rf.feature_importances_
feature_names = x.columns
feat_imp = pd.Series(importances, index=feature_names).sort_values(ascending=False)

plt.figure(figsize=(8,5))
sns.barplot(x=feat_imp, y=feat_imp.index)
plt.title("Feature Importance - Random Forest")
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.tight_layout()
plt.show()

import shap

# Create explainer using scaled data (because model was trained on it)
explainer = shap.Explainer(best_rf, x_train_scaled)

# Wrap test data in DataFrame to retain feature names
x_test_scaled_dataset = pd.DataFrame(x_test_scaled, columns=x.columns)

# Compute SHAP values (disable strict check)
shap_values = explainer(x_test_scaled_dataset, check_additivity=False)

# Global Explanation - Summary Bar Plot
shap.summary_plot(shap_values, x_test_scaled_dataset, plot_type="bar")
```

